

CS575 Project_IBM dataset

Name- Vipin Gupta

Roll- 2011MT22

▼ Downloading & Exporting the dataset

```
from pandas_datareader import data as pdr
from datetime import datetime

#download data
ibm = pdr.DataReader('IBM', 'yahoo', start=datetime(2014, 8, 1), end=datetime(2016, 11, 30))


#print first few lines of data
print(ibm.head())
```

	High	Low	...	Volume	Adj Close
Date			...		
2014-08-01	191.500000	188.860001	...	5181100.0	143.561371
2014-08-04	189.949997	188.600006	...	2125900.0	143.933304
2014-08-05	189.199997	186.440002	...	3307900.0	142.005493
2014-08-06	186.880005	184.440002	...	3847000.0	141.982544
2014-08-07	186.679993	183.580002	...	2708600.0	140.707535

[5 rows x 6 columns]

```
#export and save as csv files
ibm.to_csv('IBM_stock.csv', sep=',')
```

```
#Visulaizing the close data
import matplotlib.pyplot as plt
ibm["Close"].plot()
```



▼ Statistical analysis like ACF, PACF, ADF, KPSS Test

```
#Importing libraries
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller,kpss
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

#ADF Test

def adf_test(atr):

    #Perform Dickey-Fuller test:
    timeseries = ibm[atr].dropna()
    print ('Results of Dickey-Fuller Test for ',atr,'\n')
    dfctest = adfuller(timeseries, autolag='AIC')
    dfcoutput = pd.Series(dfctest[0:4], index=['Test Statistic','p-value','#Lags Used','Number of Observations Used'])
    for key,value in dfctest[4].items():
        dfcoutput['Critical Value (%s)'%key] = value
    print(dfcoutput)

#apply adf test on the series
```

```
adf_test('Close')
```

Results of Dickey-Fuller Test for Close

Test Statistic	-2.279273
p-value	0.178740
#Lags Used	0.000000
Number of Observations Used	588.000000
Critical Value (1%)	-3.441520
Critical Value (5%)	-2.866468
Critical Value (10%)	-2.569394

dtype: float64

The p value obtained is greater than significance level of 0.05 and test statistic is higher than any of the critical values so we cant reject the null hypothesis so the time series is non stationary.

```
#KPSS Test
```

```
def kpss_test(atr):
    timeseries = ibm[atr].dropna()
    print ('Results of KPSS Test for ',atr)
    kpsstest = kpss(timeseries, regression='c')
    kpss_output = pd.Series(kpsstest[0:3], index=['Test Statistic','p-value','Lags Used'])
    for key,value in kpsstest[3].items():
        kpss_output['Critical Value (%s)'%key] = value
    print (kpss_output)
```

```
kpss_test('Close')
```

Results of KPSS Test for Close

Test Statistic	1.268862
p-value	0.010000
Lags Used	19.000000
Critical Value (10%)	0.347000
Critical Value (5%)	0.463000
Critical Value (2.5%)	0.574000
Critical Value (1%)	0.739000


dtype: float64

/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/stattools.py:1685: FutureWarning: The behavior of using lags=Nor warn(msg, FutureWarning)


/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/stattools.py:1709: InterpolationWarning: p-value is smaller than warn("p-value is smaller than the indicated p-value", InterpolationWarning)

The p value is significant less than 0.05 hence we can reject the null hypothesis so series is non stationary

```
# ACF Test of differenced data
plot_acf(ibm['Close'].dropna(), lags=10)
plt.show()
```



```
# PACF Test of differenced data
plot_pacf(ibm['Close'].dropna(), lags=10)
plt.show()
```




```
#Differencing to make data as stationary

#Differencing the data
ibm['diff'] = ibm['Close'].diff(periods=1)

#Visulaizing the differenced data
ibm["diff"].plot()
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fc6ec551fd0>



```
# ADF Test of differenced data
adf_test('diff')
```

Results of Dickey-Fuller Test for diff

Test Statistic	-1.843371e+01
p-value	2.166547e-30
#Lags Used	1.000000e+00
Number of Observations Used	5.860000e+02
Critical Value (1%)	-3.441558e+00
Critical Value (5%)	-2.866485e+00
Critical Value (10%)	-2.569403e+00

dtype: float64

```
# KPSS Test of differenced data
kpss_test('diff')
```


Results of KPSS Test for diff

Test Statistic	0.251866
p-value	0.100000
Lags Used	19.000000
Critical Value (10%)	0.347000
Critical Value (5%)	0.463000
Critical Value (2.5%)	0.574000
Critical Value (1%)	0.739000

dtype: float64

/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/stattools.py:1685: FutureWarning: The behavior of using lags=Nor warn(msg, FutureWarning)


/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/stattools.py:1711: InterpolationWarning: p-value is greater than warn("p-value is greater than the indicated p-value", InterpolationWarning)



```
# ACF Test of differenced data
plot_acf(ibm['diff'].dropna(), lags=10)
plt.show()
```

```
Autocorrelation

# PACF Test of differenced data
plot_pacf(ibm['diff'].dropna(), lags=10)
plt.show()
```




Exponential

```
import numpy as np
from statsmodels.tsa.holtwinters import ExponentialSmoothing
```

```
n = int(len(ibm["Close"])*0.8)
data = ibm['Close'].to_numpy()
train2 = data[:n]
test2 = data[n:]
date = (ibm.index)
```


```
Exp_model = ExponentialSmoothing(ibm.Close,trend='mul',seasonal='mul',seasonal_periods=4)
ibm['Pred_Exp'] = Exp_model.fit(smoothing_level = 0.9,smoothing_slope= 0.1,smoothing_seasonal = 0.2).fittedvalues.shift(0)
```

/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/base/tsa_model.py:219: ValueWarning: A date index has been provided but is not used when e.g. forecasting., ValueWarning)



```
plt.figure(figsize=(16,8))
plt.plot(date,data, label='Original Data')
plt.plot(date[:n],ibm.Pred_Exp[:n], label='Predicted Data')
plt.plot(date[n:],ibm.Pred_Exp[n:], label='Predicted Data(20%)')
plt.legend()
plt.title('Prediction using Exponential Smoothing')
```

Text(0.5, 1.0, 'Prediction using Exponential Smoothing')



```
#Calculation of MSE for comparing the model
rmse2 = (np.mean(np.power((np.array(test2)-np.array(ibm.Pred_Exp[n:])),2)))
print('MSE value using Exponential Smoothing model: ',rmse2)
```

MSE value using Exponential Smoothing model: 3.4172334666694675

ARIMA

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from pandas.plotting import lag_plot
from pandas import datetime
from statsmodels.tsa.arima_model import ARIMA
from sklearn.metrics import mean_squared_error

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:5: FutureWarning: The pandas.datetime class is deprecated
"""

# Importing libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# Importing data
df = pd.read_csv('IBM_stock.csv')
df.head()
```

	Date	High	Low	Open	Close	Volume	Adj Close
0	2014-08-01	191.500000	188.860001	190.500000	189.149994	5181100.0	143.561371
1	2014-08-04	189.949997	188.600006	189.350006	189.639999	2125900.0	143.933304
2	2014-08-05	189.199997	186.440002	188.750000	187.100006	3307900.0	142.005493
3	2014-08-06	186.880005	184.440002	185.360001	185.970001	3847000.0	141.982544
4	2014-08-07	186.679993	183.580002	186.639999	184.300003	2708600.0	140.707535

```
# Extracting the required columns
df = df[['Date', 'Close']]
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 589 entries, 0 to 588
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Date    589 non-null    object
1   Close   589 non-null    float64
dtypes: float64(1), object(1)
memory usage: 9.3+ KB

# Changing the Date column to proper DateTime object
df.Date = pd.to_datetime(df.Date)
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 589 entries, 0 to 588
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Date    589 non-null    datetime64[ns]
1   Close   589 non-null    float64
dtypes: datetime64[ns](1), float64(1)
memory usage: 9.3 KB


# Making Date column to be the index
df.columns=['Date','Closing price']
df.set_index('Date', inplace=True)
df.head()
```

Closing price

Date


```
# Plot
df.plot(figsize=(20,10), linewidth=5,fontsize=20);
plt.xlabel('Year', fontsize=30)
```

Text(0.5, 0, 'Year')



```
# Seeing the trend more clearly
df.rolling(24).mean().plot(figsize=(20,10), linewidth=5,fontsize=20);
plt.xlabel('Year', fontsize=30)
# Overall a rise here
```


Text(0.5, 0, 'Year')



```
# We can see that there is no specific seasonality here
# Removing trend
df.diff().plot(figsize=(20,10), linewidth=5,fontsize=20);
```


```
plt.xlabel('Year', fontsize=30)
```

Text(0.5, 0, 'Year')



```
# 2nd order differencing
df.diff(periods=2).plot(figsize=(20,10), linewidth=5,fontsize=20);
plt.xlabel('Year', fontsize=30)
```

Text(0.5, 0, 'Year')




```
# Let's take a look at its auto-corelation plots
# Before that we'll have to do manual differencing
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

def difference(data, lag):
    diff= []

    for i in range(lag, len(data)):
        value = data[i] - data[i-lag]
        diff.append(value)
```

```
return pa.Series(diff)

df_close = df['Closing price']
X = df_close.values
diff = difference(X,1)
plt.plot(diff)
df_diff = pd.DataFrame(diff)
plot_acf(df_diff, lags=30)
plot_pacf(df_diff, lags=30)
```



```
# Forecasting
from sklearn.metrics import mean_squared_error
from statsmodels.tsa.arima_model import ARIMA
df = df.astype(np.float64)
Y = df.values
size = int(len(Y)*0.66)
train, test = Y[0:size], Y[size:len(Y)]
```

```
history = [x for x in train]
predictions = list()


for t in range(len(test)):
    model = ARIMA(history, order=(0,1,0))
    model_fit = model.fit(dispatch=0)
```



```
output = model_fit.forecast()
pred = output[0]
predictions.append(pred)
obs = test[t]
history.append(obs)
print('predicted=%f, expected=%f'%(pred, obs))
```

predicted=122.568396	expected=126.099998
predicted=125.937498	expected=132.449997
predicted=132.304239	expected=133.080002
predicted=132.936233	expected=133.770004
predicted=133.628367	expected=132.399994
predicted=132.255223	expected=132.800003
predicted=132.656619	expected=134.500000
predicted=134.361294	expected=132.029999
predicted=131.885391	expected=131.029999
predicted=130.883231	expected=134.369999
predicted=134.232010	expected=136.300003
predicted=136.167214	expected=137.800003
predicted=137.671306	expected=137.800003
predicted=137.671628	expected=140.149994
predicted=140.027799	expected=139.070007
predicted=138.945430	expected=140.410004
predicted=140.289061	expected=140.190002
predicted=140.068814	expected=142.360001
predicted=142.244470	expected=142.779999
predicted=142.665787	expected=142.960007
predicted=142.846518	expected=144.789993
predicted=144.681268	expected=147.039993
predicted=146.937035	expected=147.089996
predicted=146.987411	expected=148.630005
predicted=148.531416	expected=148.100006
predicted=148.000370	expected=145.399994
predicted=145.294062	expected=147.949997
predicted=147.850480	expected=148.399994
predicted=148.301801	expected=149.330002
predicted=149.234281	expected=148.410004
predicted=148.312306	expected=151.449997
predicted=151.359806	expected=152.520004
predicted=152.432582	expected=152.070007
predicted=151.981722	expected=150.000000
predicted=149.907007	expected=150.020004
predicted=149.927279	expected=148.250000
predicted=148.153310	expected=149.350006
predicted=149.256138	expected=149.250000
predicted=149.156118	expected=149.630005
predicted=149.537235	expected=151.229996
predicted=151.141190	expected=151.160004
predicted=151.071242	expected=151.720001
predicted=151.632752	expected=152.529999
predicted=152.444836	expected=144.000000
predicted=143.895244	expected=146.110001
predicted=146.010371	expected=149.300003
predicted=149.207971	expected=148.500000
predicted=148.406336	expected=148.809998
predicted=148.717262	expected=149.080002
predicted=148.988098	expected=150.470001
predicted=150.381489	expected=147.070007
predicted=146.973934	expected=145.940002
predicted=145.841574	expected=145.270004
predicted=145.170277	expected=144.130005
predicted=144.027919	expected=144.250000
predicted=144.148416	expected=146.470001
predicted=146.373658	expected=147.289993
predicted=147.195714	expected=147.339996

```
plt.plot(test)
plt.plot(predictions, color='red')
plt.show()
```



```
#Calculation of MSE for comparing the model
```

```
difference_array = np.subtract(test, predictions)
squared_array = np.square(difference_array)
mse = squared_array.mean()
mse
```

3.1982815526323543

▼ LSTM

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from keras.models import Sequential
from keras.layers import Dense,LSTM,Dropout
```

importing the training data

```
data = pd.read_csv('IBM_stock.csv')
```

choosing the close column

```
data["Close"]=pd.to_numeric(data.Close,errors='coerce') #turning the Close column to numeric
data = data.dropna() #remeving the NA values
trainData = data.iloc[:,4:5].values #selecting only the closing prices for training
```

scaling the values in the range of 0-1 for best preformances

```
sc = MinMaxScaler(feature_range=(0,1))
trainData = sc.fit_transform(trainData)
trainData.shape

(589, 1)
```

preparing the data for LSTM

since its a time series problem we took 60 as timestep for our learning : given 60 closing values as an input data the 61st value is our output

```
X_train = []
y_train = []

for i in range (60,589): #60 : timestep // 1149 : length of the data
    X_train.append(trainData[i-60:i,0])
    y_train.append(trainData[i,0])

X_train,y_train = np.array(X_train),np.array(y_train)
```

ps : LSTM take a 3D tensor (seq_len,timestep,batch_size)

```
X_train = np.reshape(X_train,(X_train.shape[0],X_train.shape[1],1)) #adding the batch_size axis
X_train.shape

(529, 60, 1)
```

building the model

```
model = Sequential()

model.add(LSTM(units=100, return_sequences = True, input_shape =(X_train.shape[1],1)))
model.add(Dropout(0.2))

model.add(LSTM(units=100, return_sequences = True))
model.add(Dropout(0.2))

model.add(LSTM(units=100, return_sequences = True))
model.add(Dropout(0.2))

model.add(LSTM(units=100, return_sequences = False))
```

```
model.add(Dropout(0.2))


model.add(Dense(units =1))
model.compile(optimizer='adam',loss="mean_squared_error")
```

```
hist = model.fit(X_train, y_train, epochs = 20, batch_size = 32, verbose=2)
```

```
Epoch 1/20
17/17 - 9s - loss: 0.0399
Epoch 2/20
17/17 - 3s - loss: 0.0095
Epoch 3/20
17/17 - 3s - loss: 0.0072
Epoch 4/20
17/17 - 3s - loss: 0.0060
Epoch 5/20
17/17 - 3s - loss: 0.0068
Epoch 6/20
17/17 - 3s - loss: 0.0079
Epoch 7/20
17/17 - 3s - loss: 0.0059
Epoch 8/20
17/17 - 3s - loss: 0.0051
Epoch 9/20
17/17 - 3s - loss: 0.0060
Epoch 10/20
17/17 - 3s - loss: 0.0048
Epoch 11/20
17/17 - 3s - loss: 0.0047
Epoch 12/20
17/17 - 3s - loss: 0.0042
Epoch 13/20
17/17 - 3s - loss: 0.0043
Epoch 14/20
17/17 - 3s - loss: 0.0046
Epoch 15/20
17/17 - 3s - loss: 0.0038
Epoch 16/20
17/17 - 3s - loss: 0.0038
Epoch 17/20
17/17 - 3s - loss: 0.0039
Epoch 18/20
17/17 - 3s - loss: 0.0036
Epoch 19/20
17/17 - 3s - loss: 0.0037
Epoch 20/20
17/17 - 3s - loss: 0.0040
```

ploting the training loss

```
plt.plot(hist.history['loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train'], loc='upper left')
plt.show()
```



testing the model on new data

```
testData = pd.read_csv('IBM_stock.csv') #importing the test data
testData["Close"]=pd.to_numeric(testData.Close,errors='coerce') #turning the close column to numerical type
testData = testData.dropna() #dropping the NA values
testData = testData.iloc[:,4:5] #selecting the closing prices for testing
y_test = testData.iloc[60:,0:].values #selecting the labels
#input array for the model
inutClosing = testData.iloc[:,0:1].values
```

```
inputClosing_scaled = sc.transform(inputClosing)
inputClosing_scaled.shape
X_test = []
length = len(testData)
timestep = 60
for i in range(timestep,length): #doing the same preivous preprocessing
    X_test.append(inputClosing_scaled[i-timestep:i,0])
X_test = np.array(X_test)
X_test = np.reshape(X_test,(X_test.shape[0],X_test.shape[1],1))
X_test.shape
```


(529, 60, 1)

```
y_pred = model.predict(X_test) #predicting the new values
```

```
predicted_price = sc.inverse_transform(y_pred) #inversing the scaling transformation for plotting
```

ploting the results

```
plt.plot(y_test, color = 'blue', label = 'Actual Stock Price')
plt.plot(predicted_price, color = 'red', label = 'Predicted Stock Price')
plt.title('GOOG stock price prediction')
plt.xlabel('Time')
plt.ylabel('Stock Price')
plt.legend()
plt.show()
```



```
#Calculation of MSE for comparing the model
difference_array = np.subtract(y_test, predicted_price)
squared_array = np.square(difference_array)
mse = squared_array.mean()
mse
```

15.992904502728457